

**THE CHAMPIONSHIP CUBE: A FRONT OFFICE FRAMEWORK FOR BUILDING NBA  
CHAMPIONS**

## Introduction

What separates the last two teams standing in the NBA Finals, and can a front office diagnose where a contender is vulnerable? Existing models rate team strength but rarely identify the weak link. We propose the Championship Cube: Team, Player and Coach Faces, plus organisational Front Office, Controllable External and Uncontrollable External Faces.

## Methods

We analyse 37 matched champion–runner-up pairs from 1989–90 to 2025–26 using Basketball-Reference, spanning the Physical ( $n=9$ ), Transition ( $n=13$ ) and Analytics ( $n=15$ ) Eras. The matched design holds the Finals stage constant and evaluates within-pair differences. Team Face is SRS, Player Face is combined Top-3 BPM, and Coach Face is regular-season wins minus Pythagorean expectation. Conditional logistic models estimate within-pair differences using odds ratios (ORs) and 95% model-based confidence intervals. A two-predictor SRS-plus-BPM model uses leave-one-pair-out validation and a temporal holdout: 22 pre-2012 pairs for training and 15 later pairs for testing. Organisational faces use a 0–2 rubric, piloted on 10 pairs by one coder using pre-Finals evidence.


## Results

Team strength and top-three player quality each separated champions from runners-up. SRS produced an OR of 1.37 (95% CI 1.08–1.75,  $p=0.011$ ) and favoured the champion in 25 of 37 pairs. Top-3 BPM produced an OR of 1.26 (95% CI 1.07–1.48,  $p=0.005$ ) and favoured the champion in 28 of 37 pairs. Their classification performance did not differ significantly (McNemar  $p=0.55$ ; within-pair  $r=0.55$ ). The two-predictor model classified 26 of 37 pairs under leave-one-pair-out validation and 11 of 15 later pairs in the temporal holdout. Regular-season coaching overperformance, payroll relative to the salary cap and draft hit rate did not distinguish champions from runners-up. Organisational scores clustered at the top of the rubric: Front Office received a score of 2 in 19 of 20 ratings, with negligible pair differences. The SRS gap was descriptively larger in later eras, but the difference was not significant (ANOVA  $p=0.43$ ).

| Face   | Measure                        | OR per unit (95% CI) | P     | Champion-favouring pairs |
|--------|--------------------------------|----------------------|-------|--------------------------|
| Team   | SRS                            | 1.37 (1.08–1.75)     | 0.011 | 25/37 (67.6%)            |
| Player | Combined Top-3 BPM             | 1.26 (1.07–1.48)     | 0.005 | 28/37 (75.7%)            |
| Coach  | Wins – Pythagorean expectation | —                    | 0.650 | 18/37 (48.6%)            |

|              |                      |   |       |               |
|--------------|----------------------|---|-------|---------------|
| Front Office | Payroll / salary cap | — | 0.330 | 23/37 (62.2%) |
| Front Office | Draft hit rate       | — | 0.980 | 16/37 (43.2%) |

Figure 1. Champion vs runner-up within each Finals (green = champion higher; red = lower)  
SRS: champion higher in 25/37 pairs; top-3 BPM: 28/37



Conclusion

Among Finals teams, roster strength—not spending, draft record or regular-season coaching overperformance—was the clearest separator. The two-predictor signal retained moderate temporal-holdout performance. The Championship Cube offers front offices a structured checklist for locating a contender’s weakest face, while KIRIN provides an exploratory six-face diagnostic. The Finals-only, range-restricted sample limits generalisation, and the organisational faces require tighter rubric anchors and independent double-coding before formal inference.