

The Championship Cube: A Six-Face Diagnostic Framework for Evaluating NBA Championship Contention

Introduction

What separates the last two teams standing in the NBA Finals, and can a front office diagnose where a contender is vulnerable? Existing models rate team strength but rarely identify the weak link. We propose the Championship Cube: Team, Player and Coach Faces, plus organisational Front Office, Controllable External and Uncontrollable External Faces.

Methods

We analyse 37 matched champion–runner-up pairs from 1989–90 to 2025–26 using Basketball-Reference. Matching holds the Finals stage constant. Team Face is SRS, Player Face is combined Top-3 BPM, Coach Face is regular-season wins minus Pythagorean expectation, and Front Office is proxied by payroll relative to the salary cap and draft hit rate. Conditional logistic models estimate odds ratios (ORs) with 95% model-based confidence intervals; franchise-clustered bootstrap intervals were wider but still excluded 1. A two-predictor SRS-plus-BPM model uses leave-one-pair-out validation and a temporal holdout: trained on 22 pre-2012 pairs, tested on 15 later pairs. Organisational faces use a 0–2 rubric, piloted on 10 pairs by one coder using pre-Finals evidence. Data, code and the KIRIN prototype are open source: github.com/ChampionshipCube-Sloan2027.

Results


Team strength and top-three player quality each separated champions from runners-up (Table 1, Figure 1). SRS produced an OR of 1.37 (95% CI 1.08–1.75, $p = 0.011$) and favoured the champion in 25 of 37 pairs. Top-3 BPM produced an OR of 1.26 (95% CI 1.07–1.48, $p = 0.005$) and favoured the champion in 28 of 37 pairs. Their classification accuracy did not differ significantly (McNemar $p = 0.55$). The two-predictor model classified 26/37 pairs under leave-one-pair-out validation and 11/15 pairs in the temporal holdout. Neither regular-season coaching overperformance nor the Front Office proxies—payroll relative to the salary cap and draft hit rate—showed statistically significant differences between champions and runners-up. Pilot organisational scores clustered at the top of the rubric (Front Office: 19 of 20 rated 2), with negligible champion–runner-up differences.

Table 1. Matched-pair results

Face	Measure	OR per unit (95% CI)	P	Champion-favouring pairs
------	---------	-------------------------	---	-----------------------------

Team	SRS	1.37 (1.08–1.75)	0.011	25/37 (67.6%)
Player	Combined Top-3 BPM	1.26 (1.07–1.48)	0.005	28/37 (75.7%)
Coach	Wins – Pythagorean expectation	—	0.650	18/37 (48.6%)
Front Office	Payroll / salary cap	—	0.330	23/37 (62.2%)
Front Office	Draft hit rate	—	0.980	16/37 (43.2%)

Figure 1. Champion vs runner-up within each Finals (green = champion higher; red = lower)
SRS: champion higher in 25/37 pairs; top-3 BPM: 28/37



Conclusion

Among Finals teams, team and player strength were the clearest separators; coaching overperformance, payroll and draft record showed no statistically significant separation. The Championship Cube offers front offices a structured checklist for locating a contender's weakest face, while KIRIN provides an exploratory six-face diagnostic. Predictors were selected on these same pairs, the Finals-only sample is range-restricted, and the organisational faces require tighter rubric anchors and independent double-coding before formal inference.